

Assignment 5

Basic CNN:

Model:

In the architecture, 2 convolutional layers are used, followed by 3 fully connected layers. Third layer provides the output. Used Relu for all the layers. Last layer output param should be 101 that's the classes. Used CrossEntropyLoss.

The following code was used:

```
21 class basic_cnn(pl.LightningModule):
22     def __init__(self):
23         # Define Model Here
24         super(basic_cnn,self).__init__()
25         #First 2d convolutional layer
26         # 1 is input channel 32 is outputting feature size, 3 is kernel size
27         self.conv1 = nn.Conv2d(3, 32, 3)
28         # Se 2d convolutional layer, takes 32 input from previous layer
29         self.conv2 = nn.Conv2d(32, 64, 3)
30
31         # First fully connected layer
32         self.fc1 = nn.Linear(57600,256)
33         # Second fully connected layer
34         self.fc2 = nn.Linear(256,128)
35         # Third fully connect layer that outputs 101 labels
36         self.out = nn.Linear(128,101)
37         self.lr = 0.01
38         self.loss = nn.CrossEntropyLoss()
39         #32x57600 and 512x256
40
41
42     def forward(self, x):
43         # Define Forward Pass Here
44         batch_size, _, _, _ = x.size()
45         #print(x.size())
46
47         x = F.relu(self.conv1(x))
48         x = F.relu(self.conv2(x))
49         x = F.max_pool2d(x, 2)
50         x = x.view(batch_size,-1)
51         x = F.relu(self.fc1(x))
52         x = F.relu(self.fc2(x))
53         return self.out(x)
```

Preparing the data:

The following code was used to prepare data:

```

6 from pl_bolts.transforms.dataset_normalizations import cifar10_normalization
7
8 class Food101(pl.LightningDataModule):
9     def prepare_data(self):
10         transform=transforms.Compose([
11             transforms.RandomResizedCrop(128),
12             transforms.RandomCrop(64),
13             transforms.RandomHorizontalFlip(),
14             transforms.ToTensor(),
15             transforms.Normalize([0.5, 0.5, 0.5],
16                                 [0.229, 0.224, 0.225]))
17         #cifar10_normalization()
18

```

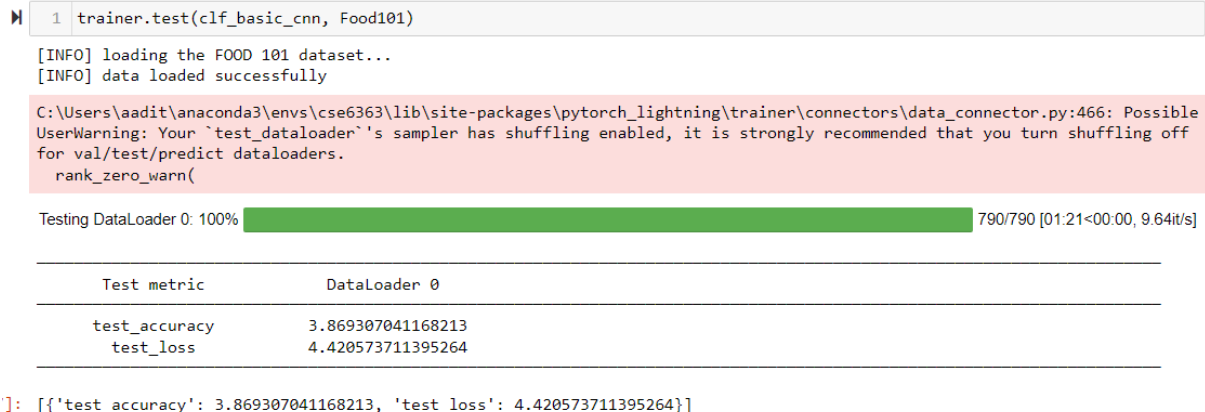
Results:

Training Loss: 4.3747

Validation Loss: 2.9199

Evaluating test data:

Below image shows accuracy and loss results for basic cnn on testing set:



All Convolutional Net:

Model:

In the architecture, 3 convolutional layers are used, followed by a pooling layer. Flatten layer provides the output. Used Relu for all the conv2d layers. Last layer output param should be 101 that's the classes. Used CrossEntropyLoss.

The following code was used:

```

17 class all_conv_net(pl.LightningModule):
18     def __init__(self):
19         # Define Model Here
20         super(all_conv_net,self).__init__()
21         #First 2d convolutional layer
22         # 1 is input channel 32 is outputting feature size, 3 is kernel size
23         self.conv1 = nn.Conv2d(3, 32, 3)
24         # Se 2d convolutional layer, takes 32 input from previous layer
25         self.conv2 = nn.Conv2d(32, 64, 3)
26
27         self.conv3 = nn.Conv2d(64, 101, 3)
28         # First fully connected layer
29         #self.fc1 = nn.Linear(57600,256)
30         # Second fully connected layer
31         #self.fc2 = nn.Linear(256,128)
32         # Third fully connect layer that outputs 101 labels
33         #self.out = nn.Linear(128,101)
34         self.lr = 0.01
35         self.loss = nn.CrossEntropyLoss()
36         #32x57600 and 512x256
37
38
39
40     def forward(self, x):
41         # Define Forward Pass Here
42         batch_size, _, _, _ = x.size()
43         #print(x.size())
44
45         x = F.relu(self.conv1(x))
46         x = F.relu(self.conv2(x))
47         x = F.relu(self.conv3(x))
48         x = F.max_pool2d(x, 2)
49         #x = x.view(batch_size,-1)
50         #x = F.relu(self.fc1(x))
51         #x = F.relu(self.fc2(x))
52         x = x.flatten(1)
53         return x
54

```

Preparing the data:

The same data is used from the basic CNN model.

Results:

Training Loss: 6.7058

Validation Loss: 6.6397

Evaluating test data:

Below image shows accuracy and loss results for all convolutional network on testing set:

```
7]: 1 trainer.test(clf_all_conv_net, Food101)

[INFO] loading the FOOD 101 dataset...
[INFO] data loaded successfully

Testing DataLoader 0: 100%  790/790 [02:29<00:00, 5.27it/s]



| Test metric   | DataLoader 0        |
|---------------|---------------------|
| test_accuracy | 0.23366336524486542 |
| test_loss     | 6.738712787628174   |



it[17]: [{'test accuracy': 0.23366336524486542, 'test loss': 6.738712787628174}]
```

Regularization:

In the architecture, 2 convolutional layers are used, followed by 3 fully connected layers. Third layer provides the output. Used Relu for all the layers. Added 2 dropout layers with 25% dropout. Last layer output param should be 101 that's the classes. Used CrossEntropyLoss.

Model:The following code was used:

```

16 class reg(pl.LightningModule):
17     def __init__(self):
18         # Define Model Here
19         super(reg,self).__init__()
20         #First 2d convolutional layer
21         # 1 is input channel 32 is outputting feature size, 3 is kernel size
22         self.conv1 = nn.Conv2d(3, 32, 3)
23         # Se 2d convolutional layer, takes 32 input from previous layer
24         self.conv2 = nn.Conv2d(32, 64, 3)
25         # First fully connected layer
26         self.fc1 = nn.Linear(57600,256)
27         # Second fully connected layer
28         self.fc2 = nn.Linear(256,128)
29         # Third fully connect layer that outputs 101 labels
30         self.out = nn.Linear(128,101)
31
32         #Define proportion or neurons to dropout
33         self.dropout = nn.Dropout(0.25)
34
35
36         self.lr = 0.01
37         self.loss = nn.CrossEntropyLoss()
38         #32x57600 and 512x256
39
40
41     def forward(self, x):
42         # Define Forward Pass Here
43         batch_size, _, _, _ = x.size()
44         #print(x.size())
45
46         x = F.relu(self.conv1(x))
47         x = F.relu(self.conv2(x))
48         x = F.max_pool2d(x, 2)
49         x = x.view(batch_size,-1)
50         x = self.dropout(x)
51         x = F.relu(self.fc1(x))
52         x = self.dropout(x)
53         x = F.relu(self.fc2(x))
54         return self.out(x)

```

Preparing the data:

Applied RandAugment function to augment the data. Added 2 dropout layers with 25% dropout. The following code was used to prepare data:

```

7
8 class aug_Food101(pl.LightningDataModule):
9     def prepare_data(self):
10         transform=transforms.Compose([
11             transforms.RandomResizedCrop(128),
12             transforms.RandAugment(),
13             transforms.RandomCrop(64),
14             transforms.RandomHorizontalFlip(),
15             transforms.ToTensor(),
16             transforms.Normalize([0.5, 0.5, 0.5],
17                                 [0.229, 0.224, 0.225]))
18         #cifar10_normalization()
19

```

Results:

Training Loss: 4.6841

Aaditya Damle


1001955625

Validation Loss: 4.5177

Evaluating test data:

Below image shows accuracy and loss results after applying regularization to basic cnn on testing set:

```
1 trainer.test(clf_reg, aug_Food101)
[INFO] loading the FOOD 101 dataset...
[INFO] data loaded successfully

Testing DataLoader 0: 100%  790/790 [01:02:00.00, 12.55it/s]
```

Test metric	DataLoader 0
test_accuracy	2.8039603233337402
test_loss	4.498032569885254

```
[22]: [{'test_accuracy': 2.8039603233337402, 'test_loss': 4.498032569885254}]
```

Transfer Learning:

Model:

Used the pre-trained Resnet18 model. Resnet 50 is a very good model for image classification but as the model is very big takes a lot of time and resources to train the backbone. So I took the resnet18 which has less number of layers which decreased training time by good extent. The following code was used to get the pre-trained model.

```
# init a pretrained resnet
backbone = models.resnet18(pretrained=True)
num_filters = backbone.fc.in_features
layers = list(backbone.children())[:-1]
self.feature_extractor = nn.Sequential(*layers)

# use the pretrained model to classify food101
num_target_classes = 101
self.classifier = nn.Linear(num_filters, num_target_classes)

def forward(self, x):
    self.feature_extractor.eval()
    with torch.no_grad():
        representations = self.feature_extractor(x).flatten(1)
    x = self.classifier(representations)

    return x
```

Preparing the data:

Changed the input to RandomCrop function in transforms to 56. The following code was used to prepare data:

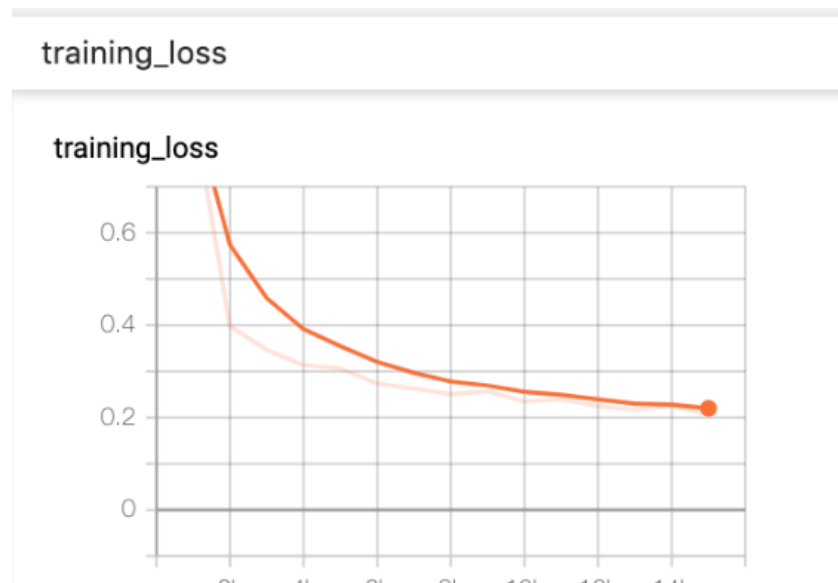
```

30         self.batch_size = batch_size
31         self.lr = lr
32
33         # Data preparation
34         dataset = Food101(data_path, transform=transforms.Compose([
35             transforms.RandomAugment(),
36             transforms.RandomCrop(56),
37             transforms.ToTensor(),
38         ]), download=True)
39
40         dataset_size = len(dataset)
41         train_size = int(dataset_size * .95)
42         val_size = dataset_size - train_size
43
44         self.train_dataset, self.val_dataset = torch.utils.data.random_split(dataset, [train_size, val_size])
45
46         # Loss function
47         self.loss_fn = nn.CrossEntropyLoss()
48

```

Results:

Training Loss:



Validation Loss: 4.43

Evaluating test data:

Below image shows accuracy and loss results for resnet18 model used for transfer learning on testing set:

1
trainer.test(model)

Testing DataLoader 0: 100%
119/119 [00:11<00:00, 10.69it/s]

Test metric	DataLoader 0
test_accuracy	4.963041305541992
test_loss	4.55490255355835

66]: [{'test_accuracy': 4.963041305541992, 'test_loss': 4.55490255355835}]

Bonus:**Model:**

Used only 1 conv2d layer in the model. The following code was used for defining model:

```
class bonus(pl.LightningModule):
    def __init__(self):
        # Define Model Here
        super(bonus,self).__init__()
        #First 2d convolutional layer
        # 1 is input channel 32 is outputting feature size, 3 is kernel size
        #self.conv1 = nn.Conv2d(3, 32, 3)
        # Se 2d convolutional layer, takes 32 input from previous layer
        #self.conv2 = nn.Conv2d(32, 64, 3)

        self.conv1 = nn.Conv2d(3, 101, 3)
        # First fully connected layer
        #self.fc1 = nn.Linear(57600,256)
        # Second fully connected layer
        #self.fc2 = nn.Linear(256,128)
        # Third fully connect layer that outputs 101 labels
        #self.out = nn.Linear(128,101)
        self.lr = 0.01
        self.loss = nn.CrossEntropyLoss()
        #32x57600 and 512x256

    def forward(self, x):
        # Define Forward Pass Here
        batch_size, _, _, _ = x.size()
        #print(x.size())

        x = F.relu(self.conv1(x))
        #x = F.relu(self.conv2(x))
        #x = F.relu(self.conv3(x))
        x = F.max_pool2d(x, 2)
        #x = x.view(batch_size, -1)
        #x = F.relu(self.fc1(x))
        #x = F.relu(self.fc2(x))
        x = x.flatten(1)
        return x
```

Preparing the data:

Changed the input to RandomCrop function in transforms to 56. The following code was used to prepare data:

